

Pointer Analysis – Part II

CS 6340


Unification vs. Inclusion

- Earlier scalable pointer analysis was context-insensitive unification-based [Steensgaard '96]
 - Pointers are either un-aliased or point to the same set of objects
 - Near-linear, but very imprecise
- Inclusion-based pointer analysis
 - Can point to overlapping sets of objects
 - Closure calculation is $O(n^3)$
 - Various optimizations [Fahndrich, Su, Heintze,...]
 - BDD formulation, simple, scalable [Berndl, Zhu]

1

Context Sensitivity


- Context sensitivity is important for precision
 - Unrealizable paths



2

Context Sensitivity

- Context sensitivity is important for precision
 - Unrealizable paths
 - Conceptually give each caller its own copy



3

Summary-Based Analysis

- Popular method for context sensitivity
- Two kinds:
 - Bottom-up
 - Top-down
- Problems:
 - Difficult to summarize pointer analysis
 - Summary-based analysis using BDD: not shown to scale [Zhu'02]


4

Cloning-Based Analysis

- Simple brute force technique
 - Clone every path through the call graph
 - Run context-insensitive algorithm on the expanded call graph
- The catch: exponential blowup

5

Cloning is exponential!




6

Recursion

- Actually, cloning is unbounded in the presence of recursive cycles
- Solution: treat all methods in a strongly-connected component as a single node


7

Recursion



8

Top 20 Sourceforge Java Apps



9

Cloning is infeasible (?)

- Typical large program has $\sim 10^{14}$ paths
- If you need 1 byte to represent a clone, would require 256 terabytes of storage
 - Registered ECC 1GB DIMMs: \$98.6 million
 - Power: 96.4 kilowatts = Power for 128 homes
 - 300 GB hard disks: $939 \times \$250 = \$234,750$
 - Time to read (sequential): 70.8 days
- Seems unreasonable!

10

BDD comes to the rescue

- There are many similarities across contexts
 - Many copies of nearly-identical results
- BDDs can represent large sets of redundant data efficiently
 - Need a BDD encoding that exploits the similarities

11

Contribution (1)

- Can represent context-sensitive call graph efficiently with BDDs and a clever context numbering scheme
 - Inclusion-based pointer analysis
 - 10^{14} contexts, 19 minutes
 - Generates all answers

12

Contribution (2)

BDD hacking is complicated →
bddbdb (BDD-based deductive database)

- Pointer analysis in 6 lines of Datalog
- Automatic translation into efficient BDD implementation
- 10x performance over hand-tuned solver (2164 lines of Java)

13


Contribution (3)

- bddbdb: General Datalog solver
 - Supports simple declarative queries
 - Easy use of context-sensitive pointer results
- Simple context-sensitive analyses:
 - Escape analysis
 - Type refinement
 - Side effect analysis
 - Many more presented in the paper

14

Call Graph Relation

- Call graph expressed as a relation
 - Five edges:
 - Calls(A,B)
 - Calls(A,C)
 - Calls(A,D)
 - Calls(B,D)
 - Calls(C,D)




15

Call Graph Relation

x_1	x_2	x_3	x_4	f
0	0	0	0	0
0	0	0	1	1
0	0	1	0	1
0	0	1	1	1
0	1	0	0	0
0	1	0	1	0
0	1	1	0	0
0	1	1	1	1
1	0	0	0	0
1	0	0	1	0
1	0	1	0	0
1	0	1	1	1
1	1	0	0	0
1	1	0	1	0
1	1	1	0	0
1	1	1	1	0


- Relation expressed as a binary function.
 - A=00, B=01, C=10, D=11



16

Binary Decision Diagrams


- Graphical encoding of a truth table



17

Binary Decision Diagrams


- Collapse redundant nodes



18

Binary Decision Diagrams


- Collapse redundant nodes



19

Binary Decision Diagrams


- Collapse redundant nodes



20

Binary Decision Diagrams


- Collapse redundant nodes



21

Binary Decision Diagrams


- Eliminate unnecessary nodes



22

Binary Decision Diagrams

- Eliminate unnecessary nodes




23

Binary Decision Diagrams

- Size is correlated to amount of redundancy, NOT size of relation
 - As the set gets larger, the number of don't-care bits *increases*, leading to fewer necessary nodes


24

Expanded Call Graph



25

Numbering Clones



26


Pointer Analysis Example



$h_1: v_1 = \text{new Object}();$
 $h_2: v_2 = \text{new Object}();$
 $v_1.f = v_2;$
 $v_3 = v_1.f;$

Input Relations

$\text{vPointsTo}(v_1, h_1)$
 $\text{vPointsTo}(v_2, h_2)$
 $\text{Store}(v_1, f, v_2)$
 $\text{Load}(v_1, f, v_3)$



Output Relations


$\text{hPointsTo}(h_1, f, h_2)$
 $\text{vPointsTo}(v_3, h_2)$

27

Inference Rule in Datalog

Heap Writes:

$\text{hPointsTo}(h_1, f, h_2) \text{ :- } \text{Store}(v_1, f, v_2),$
 $\text{vPointsTo}(v_1, h_1),$
 $\text{vPointsTo}(v_2, h_2).$



28

Context-sensitive pointer analysis

- Compute call graph with context-insensitive pointer analysis
 - Datalog rules for:
 - assignments, loads, stores
 - discover call targets, bind parameters
 - type filtering
 - Apply rules until fix-point reached
- Compute expanded call graph relation
- Apply context-insensitive algorithm to the expanded call graph

29

Datalog

- Declarative logic programming language designed for databases
 - Horn clauses
 - Operates on relations
- Datalog is expressive
 - Relational algebra:
 - Explicitly specify relational join, project, rename
 - Relational calculus:
 - Specify relations between variables; operations are implicit
 - Datalog:
 - Allows recursively-defined relations

30

Datalog → BDD

- Join, project, rename are directly mapped to built-in BDD operations
- Automatically optimizes:
 - Rule application order
 - Incrementalization
 - Variable ordering
 - BDD parameter tuning
 - Many more...

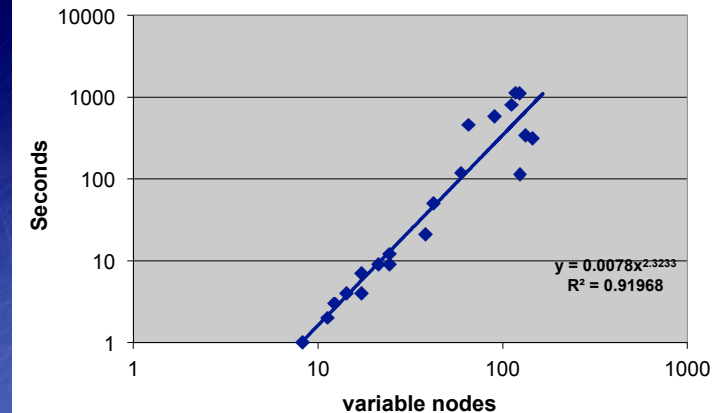
31

Experimental Results

- Top 20 Java projects on SourceForge
 - Real programs with 100K+ users each
- Using automatic bddbddb solver
 - Each analysis only a few lines of code
 - Easy to try new algorithms, new queries
- Test system:
 - Pentium 4 2.2GHz, 1GB RAM
 - RedHat Fedora Core 1, JDK 1.4.2_04, javabdd library, Joeq compiler

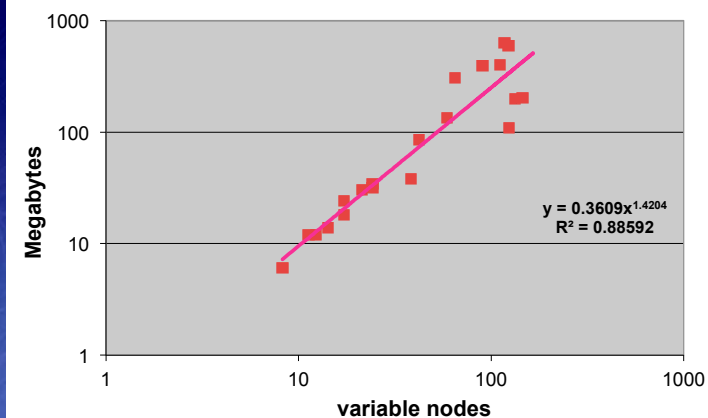
32

Analysis Time



33

Analysis memory



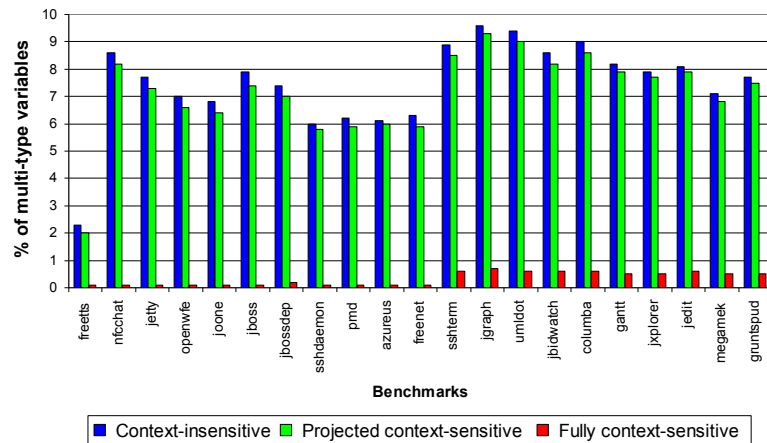
34

Multi-type variables

- A variable is multi-type if it can point to objects of different types
 - Measure of analysis precision
 - One line in Datalog
- Two ways of handling context sensitivity:
 - Projected: Merge all contexts together
 - Full: Keep each context separate

35

Comparison of Accuracy (smaller bars are better)



36

Conclusion

- The first scalable context-sensitive inclusion-based pointer analysis
 - Achieves context sensitivity by cloning
- bddbddb: Datalog → efficient BDD
- Easy to query results, develop new analyses
- Very efficient!
 - < 19 minutes, < 600mb on largest benchmark
- Complete system is publicly available at:
<http://suif.stanford.edu/bddbddb>

37